




# Multithreading

---

CrowdStrike HEROES - Cloud Workshop

Date	Topic
July 25 - 16:00	Intro to golang
July 26 - 16:00	Intro to golang (continuation)
<b>July 27 - 16:00</b>	<b>Multithreading</b>
July 28 - 16:00	Rest API
July 29 - 16:00	Unit testing, logging and monitoring
August 1 - 16:00	Workshop and Q&A
August 2 - 16:00	Deployments/Docker
August 3 - 16:00	Databases
August 4 - 16:00	Databases extended
August 5 - 13:00	Microservices contest (4h with Awards)



## Basics

Why do we need multithreading?

What is a thread?

What is a race condition?


How to fix a race condition?



Our World  
in Data

## Transistor count

50,000,000,000



OurWorldinData.org – Research and data to make progress against the world's largest problems

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser

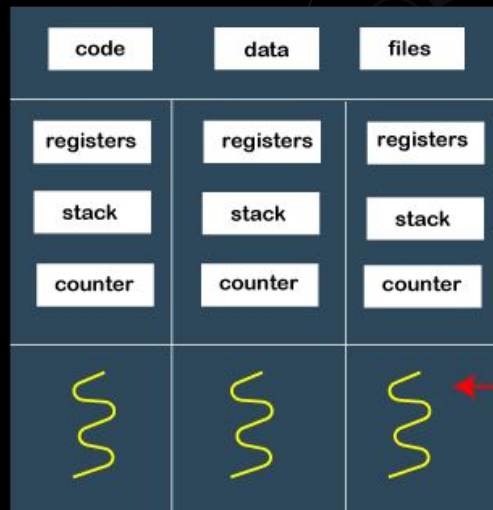
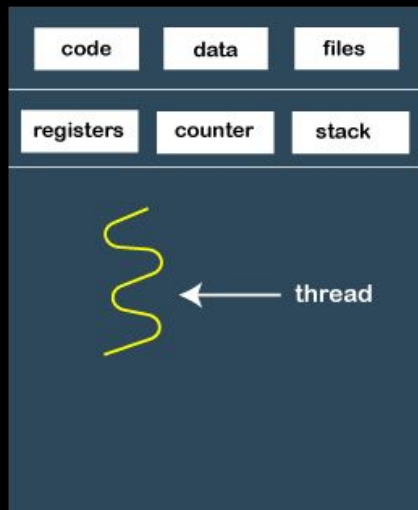
- Moore's law will cease to apply at some point
- Solution: multiple processors
- Problem: working with multiple processors requires more knowledge and skill to do it correctly



# Threads

A thread is a path of execution within a process.

- Threads within the same process run in a shared memory space



More about how the OS handles threads [here](https://static.javatpoint.com/difference/images/process-vs-thread3.png)



# Race conditions

## Concurrent updates

Thread A

```
count = count + 1
```

Thread B

```
count = count + 1
```

What is the value of count at the end, assuming `count=0` ?



## Race conditions

Concurrent updates behind the scenes

Thread A

```
temp = count  
count = temp + 1
```

Thread B

```
temp = count  
count = temp + 1
```

The value of count can be either one of these:

```
count = 2
```

```
count = 1
```

Now imagine the same problem with 100 threads :D

In machine operations the increment is split in two operations.

Some computers do actually provide an increment instruction that cannot be interrupted.

An operation that cannot be interrupted is called **atomic**.





## Solution: mutex

Mutex = mutual exclusion

Usually the operations we need to perform on a mutex are:

- `mutex.Lock()`
- `mutex.Unlock()`

These operations are usually surrounding a critical section.

Be aware that any lock waits for the other threads to unlock the mutex.

Misuse of mutexes can lead to problems: see [deadlocks](#).





# Go syntax for multithreading

Goroutines

Channels

Sync package: mutex, waitGroup



# Goroutines

- A goroutine is a lightweight thread managed by the Go runtime.

```
go f(x, y, z)
```

starts a new goroutine running

```
f(x, y, z)
```

- `f`, `x`, `y`, and `z` are **evaluated** in the current goroutine
- `f` is **executed** in the new goroutine
- Goroutines run in the same address space => access to shared memory must be synchronized

```
func say(s string) {  
    for i := 0; i < 5; i++ {  
        time.Sleep(100 *  
time.Millisecond)  
        fmt.Println(s)  
    }  
}
```

```
func main() {  
    go say("world")  
    say("hello")  
}
```

## Function closures

- We can define anonymous functions
  - Much like lambdas in other languages

```
f := func (args) retType {  
    // do stuff  
}  
ret := f(arg1, arg2)
```

- We can call it immediately by adding the arguments after the definition  

```
res := func() int {return 0} ()
```
- As you can see we can have functions with no args or no return value



## Variables in closures

### Capturing outside variables

```
v := 3
f := func () int {
    return v
}
v = 4
fmt.Println(f())
> 4
```

Using values from surrounding context uses the value at the time of the call. All further changes to the objects can be seen in the function

### Capturing variables by arguments

```
v := 3
res := func (v int) int {
    return v
} (v)
v = 4
fmt.Println(res)
> 3
```

Sending values as arguments takes the value at the time of the definition



# Channels

- Channels help us send and receive values
- The channel operator is <-

```
ch <- v    // Send v to channel ch.  
v := <-ch  // Receive from ch, and  
           // assign value to v.
```

(The data flows in the direction of the arrow.)

- Channels must be created before use:  

```
ch := make(chan int)
```
- Sends and receives block until the other side is ready

```
func sum(s []int, c chan int) {  
    sum := 0  
    for _, v := range s {  
        sum += v  
    }  
    c <- sum // send sum to c  
}  
  
func main() {  
    s := []int{7, 2, 8, -9, 4, 0}  
  
    c := make(chan int)  
    go sum(s[:len(s)/2], c)  
    go sum(s[len(s)/2:], c)  
    x, y := <-c, <-c // receive from c  
  
    fmt.Println(x, y, x+y)  
}
```

## Buffered Channels

- Provide the buffer length as the second argument to make to initialize a **buffered** channel:  
`ch := make(chan int, 100)`
- **Sends** to a buffered channel block only when the buffer is full.
- **Receives** block when the buffer is empty.

```
func main() {  
    ch := make(chan int, 2)  
    ch <- 1  
    ch <- 2  
    fmt.Println(<-ch)  
    fmt.Println(<-ch)  
}
```

## Range & Close

- A sender can close a channel to indicate that no more values will be sent.
- Receivers can test whether a channel has been closed:

```
v, ok := <-ch
```

- If `ok` is `false` there are no more values to receive and the channel is closed.
- To receive values from the channel repeatedly until it is closed:

```
for i := range c
```

```
func fibonacci(n int, c chan int) {  
    x, y := 0, 1  
    for i := 0; i < n; i++ {  
        c <- x  
        x, y = y, x+y  
    }  
    close(c)  
}
```

```
func main() {  
    c := make(chan int, 10)  
    go fibonacci(cap(c), c)  
    for i := range c {  
        fmt.Println(i)  
    }  
}
```



## Select

- The `select` statement lets a goroutine wait on multiple communication operations.
- A `select` blocks until one of its cases can run, then it executes that case.
- It chooses one at random if multiple are ready.

```
func fibonacci(c, quit chan int) {  
    x, y := 0, 1  
    for {  
        select {  
        case c <- x:  
            x, y = y, x+y  
        case <-quit:  
            fmt.Println("quit")  
            return  
        }  
    }  
}
```

## Default Selection

- The default case in a select is run if no other case is ready
- Use a default case to try a send or receive without blocking:

```
select {  
  case i := <-c:  
    // use i  
  default:  
    // receiving from c would block  
}
```



## sync.Mutex

- Go's standard library provides mutual exclusion with `sync.Mutex` and its two methods:
  - Lock
  - Unlock
- We can define a block of code to be executed in mutual exclusion by **surrounding** it with a call to `Lock` and `Unlock`

```
func (c *SafeCounter) Inc(key string) {  
    c.mu.Lock()  
    // Lock so only one goroutine at a time can access the map c.v.  
    c.v[key]++  
    c.mu.Unlock()  
}
```

## Thread pools - Wait Groups

Since it's not very productive to start too many threads we can have only a few worker goroutines and send them the tasks.

To wait for multiple goroutines to finish, we can use a *wait group*.

- Wait groups are already initialized
- Add a goroutine: `wg.Add(1)`
- Mark the end of a goroutine:  
`wg.Done()`
- Wait for all goroutines to finish: Add a goroutine: `wg.Add(1)`

```
var wg sync.WaitGroup
for i := 1; i <= 5; i++ {
    wg.Add(1)
    i := i
    go func() {
        defer wg.Done()
        worker(i)
    }()
}
wg.Wait()
```



Any  
questions?

