


# Databases

---

CrowdStrike HEROES - Cloud Workshop



Date	Topic
July 25 - 16:00	Intro to golang
July 26 - 16:00	Intro to golang (continuation)
July 27 - 16:00	Multithreading
July 28 - 16:00	Rest API
July 29 - 16:00	Unit testing, logging and monitoring
August 1 - 16:00	Workshop and Q&A
August 2 - 16:00	Deployments/Docker
<b>August 3 - 16:00</b>	<b>Databases</b>
August 4 - 16:00	Databases extended
August 5 - 13:00	Microservices contest (4h with Awards)



# What we're talking about here today

## Intro

The why and the how

## Environment setup

Docker it up, go it up, and  
Postgres the hell out of it

## CRUD!

What a database is, and how  
the basic operations look like

## Join the club

Constraints, joins, migrations

## Stretching the cord

Bulk inserts at scale, and how  
to optimize for large data  
volumes

## ACID

Because this is a fun  
workshop





## A database is where you keep your data

A database is an organized collection of structured information, or data, typically stored electronically in a computer system.

## Relational databases

Relational databases became dominant in the 1980s. Items in a relational database are organized as a set of tables with columns and rows.




## Key-Value




## Graph DB



## Column Family



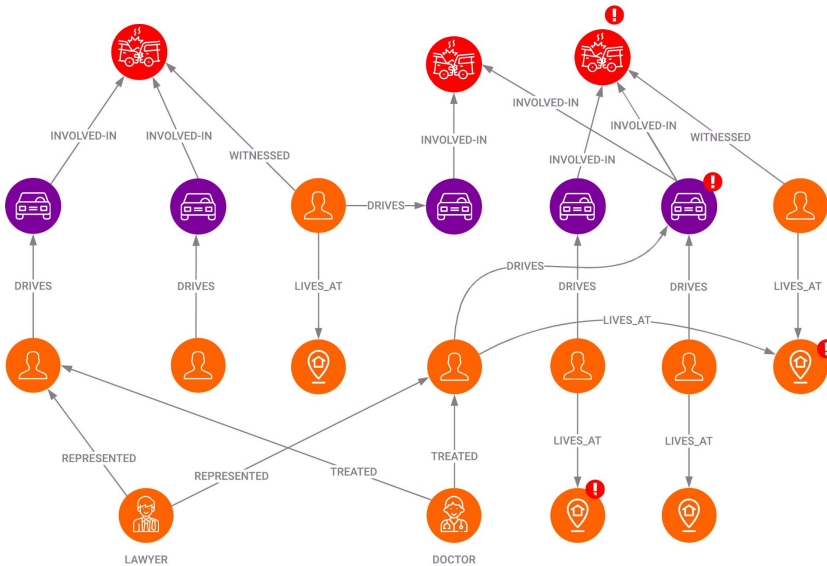
## Document



## NoSQL databases

A NoSQL, or nonrelational database, allows unstructured and semistructured data to be stored and manipulated






## Graph databases

A graph database stores data in terms of entities and the relationships between entities.



## Most Popular Database Technologies in Stack Overflow 2019 Survey



## Intro

The one and only: PostgreSQL

- Free, open source, relational database
- Older than you: Initial release was on July 8th 1996
- Very popular





# Environment setup

Docker, Golang, Postgres



# CRUD!


Create, Read, Update, Delete operations




Join the club

Constraints, joins, and migrations







SELECT \* FROM a  
INNER JOIN b ON a.key = b.key



SELECT \* FROM a  
LEFT JOIN b ON a.key = b.key



SELECT \* FROM a  
RIGHT JOIN b ON a.key = b.key




## POSTGRESQL JOINS

SELECT \* FROM a  
LEFT JOIN b ON a.key = b.key  
WHERE b.key IS NULL



SELECT \* FROM a  
RIGHT JOIN b ON a.key = b.key  
WHERE a.key IS NULL



SELECT \* FROM a  
FULL JOIN b ON a.key = b.key



SELECT \* FROM a  
FULL JOIN b ON a.key = b.key  
WHERE a.key IS NULL OR b.key IS NULL



# Stretching the cord

Bulk inserts, Indexes, B-tree



# ACID

Atomicity, Consistency, Isolation, Durability



### Atomicity

Commits finish an entire operation successfully or the database rolls back to its prior state

### Consistency

Any change maintains data integrity or is cancelled completely

### Isolation

Any read or write will not be impacted by other reads or writes of separate transactions

### Durability

Successful commits will survive permanently





# What we talked about today

## Intro

The why and the how

## Environment setup

Docker it up, go it up, and  
Postgres the hell out of it

## CRUD!

What a database is, and how  
the basic operations look like

## Join the club

Constraints, joins, migrations

## Stretching the cord

Bulk inserts at scale, and how  
to optimize for large data  
volumes

## ACID

Because this is a fun  
workshop



Any  
questions?

